

CSE 601: Distributed Systems

Toukir Ahammed

Architectural Styles

Architectural Styles

Important styles of architecture for distributed systems:

- Layered architectures
- Object-based architectures
- Data-centered architectures
- Event-based architectures

Architectural Styles

Resource-based Architectures

Resource-based Architectures

- View a distributed system as a huge collection of resources that are individually managed by components
- Resources may be *added* or *removed* by (remote) applications, and likewise can be *retrieved* or *modified*
- This approach has now been widely adopted for the Web and is known as **Representational State Transfer (REST)**
- REST is an architectural style based on transferring representations of resources from a server to a client

Resource-based Architectures

There are four key characteristics of RESTful architectures:

1. Resources are identified through a single naming scheme
2. All services offer the same interface, consisting of at most four operations
3. Messages sent to or from a service are fully self-described
4. After executing an operation at a service, that component forgets everything about the caller (stateless execution)


Resource-based Architectures

Basic Operations


Operation	Description
PUT	Create a new resource
GET	Retrieve the state of a resource in some representation
DELETE	Delete a resource
POST	Modify a resource by transferring a new state

Resource-based Architectures

Basic Operations



(a) General resource actions



(b) Web resources

Resource-based Architectures

Example: Amazon's Simple Storage Service (Amazon S3)

- *Objects* (i.e., files) are placed into *buckets* (i.e., directories)
- *Buckets* cannot be placed into buckets
- Operations on *ObjectName* in bucket *BucketName* require the following identifier:
<https://s3.amazonaws.com/BucketName/ObjectName>

Typical operations

All operations are carried out by sending HTTP requests:

- Create a bucket/object: POST, along with the URI
- Listing objects: GET on a bucket name
- Reading an object: GET on a full URI

Architectural Styles

Event based architectures

Event based Architectures

- As systems continue to grow and processes can more easily join or leave, it becomes important to have an architecture in which dependencies between processes become as loose as possible
- An architecture is needed in which there is a strong separation between *processing* and *coordination*
- *Coordination* can be discussed from two dimensions: temporal and referential.

Event based Architectures

- **Temporal coupling** means that processes that are communicating will both have to be up and running.
- **Referential coupling** generally appears in the form of explicit referencing in communication.
 - For example, a process can communicate only if it knows the name or identifier of the other processes it wants to exchange information with.

Event based Architectures

	Temporally coupled	Temporally decoupled
Referentially coupled	Direct	Mailbox
Referentially decoupled	Event-based	Shared data space

Event based Architectures

- **Direct coordination:** when processes are temporally and referentially coupled, coordination takes place directly.
 - In real life, talking over cell phones (and assuming that a cell phone has only one owner), is an example of direct communication.
- **Mailbox coordination:** when processes are temporally decoupled, but referentially coupled
 - there is no need for two communicating processes to be active at the same time to let communication take place.
 - Instead, communication takes place by putting messages in a (possibly shared) mailbox

Event based Architectures

- **Event-based coordination:** The combination of referentially decoupled and temporally coupled systems
 - processes do not know each other explicitly
 - a process **publish** a notification describing the occurrence of an event
 - processes may **subscribe** to a specific kind of notification
 - it is generally required that the subscriber is up-and-running at the time the
 - notification was published

Event based Architectures

- **Shared data space:** the combination of referentially and temporally decoupled processes
 - The key idea is that processes communicate entirely through tuples, which are structured data records consisting of several fields, very similar to a row in a database table
 - Processes can put any type of tuple into the shared data space.
 - To retrieve a tuple, a process provides a search pattern that is matched against the tuples. Any tuple that matches is returned
 - A process subscribes to certain tuples by providing a search pattern; when a process inserts a tuple into the data space, matching subscribers are notified

Publish-subscribe architectures

